

AI Persona Evals Report (Voice + Chat)

Generated: 2026-06-05T10:55:23.466112

Executive Summary

Voice-first phone agent and chat agent use the same RAG knowledge base (resume + GitHub corpus). This run generated chat groundedness metrics and retrieval metrics; voice booking metrics require real call logging (evals/call_logs.jsonl).

Part A: Voice Quality

No voice call logs found (evals/call_logs.jsonl). Run test calls via the Vapi phone number to populate logs.

Part B: Chat Groundedness

Total questions: 7. Hallucination rate: 0.0%. Accuracy (answer or correct rejection): 71.4%.

Retrieval Quality

This repo's run_evals.py prints retrieval precision/recall to stdout. If you want these persisted into metrics.json, update scripts/run_evals.py accordingly.

Failure Modes Found (Template)

1) If calendar tooling is missing/unauthenticated, booking confirmations can fall back or fail. Fix: ensure CALCOM_API_KEY/Google auth are configured and webhook parsing is strict. 2) If retrieval returns low-confidence docs, the persona should say it lacks the information and ask during the interview. Fix: tighten groundedness prompt + retrieval confidence checks. 3) Prompt injection attempts should not override system instructions. Fix: detect injection patterns and force a grounded refusal.

Conscious Tradeoff

Speed vs perfect accuracy: the system streams chat responses and keeps temperature low to reduce hallucination risk while maintaining low latency for screening UX.

What I'd Build in 2 More Weeks

- Automated end-to-end test harness that drives voice and chat scenarios and measures booking success.
- Persistent retrieval precision/recall and a judge-based hallucination grader with labeled evidence.
- Stronger structured extraction for booking (datetime/name/email) with validation + retry loops.